

SQL Moderation Hack

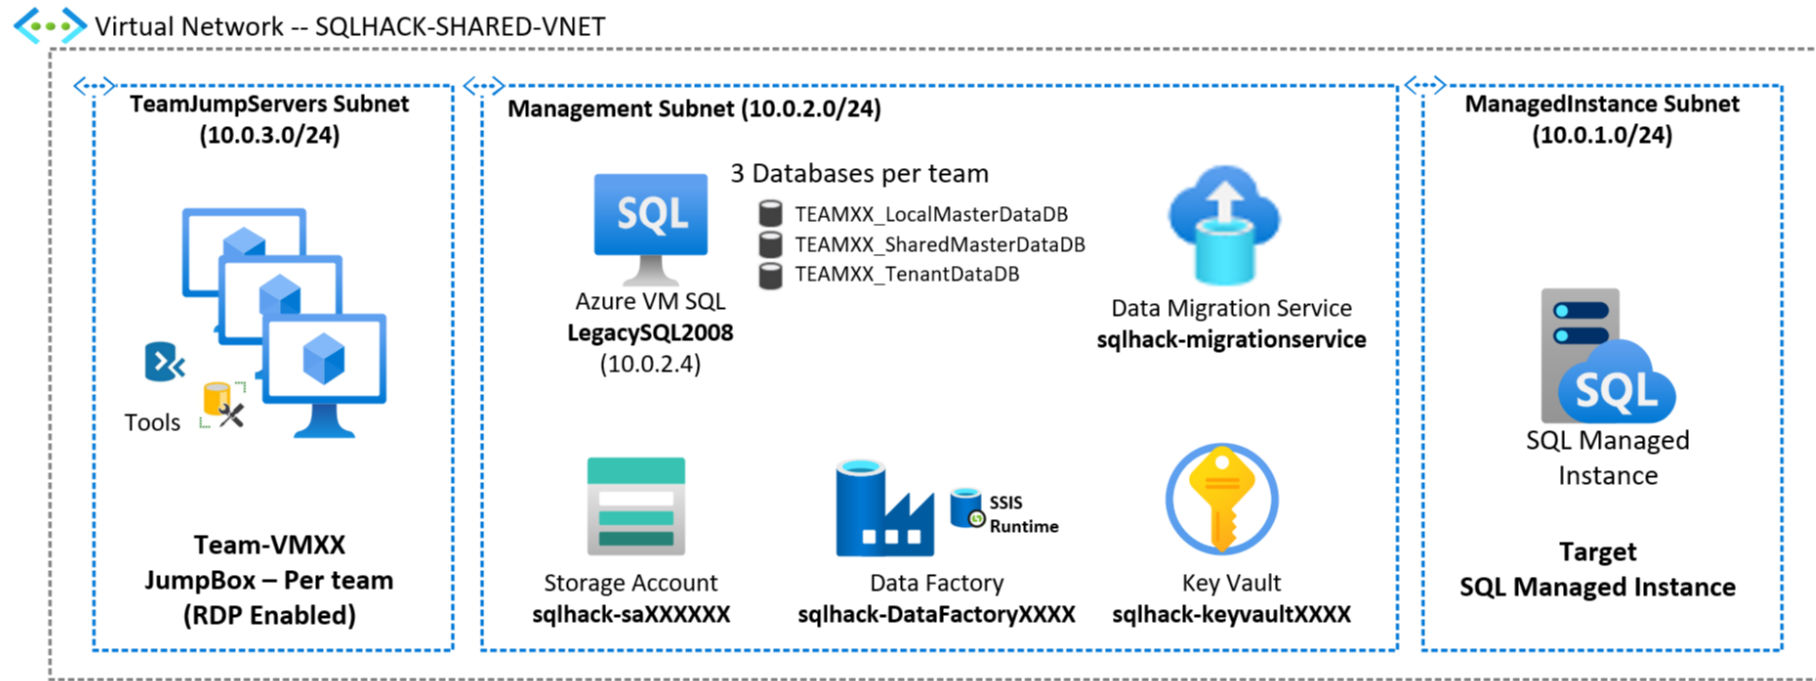
Database Administering and Monitoring Labs Step-by-step

V3.0

Contents

Migration architecture and Azure components	2
Lab Overview & Background	3
Background	3
LAB 1: Identifying performance issues, their root causes and fixing the problem	4
1. Using DMVs and the Query Store to identify performance bottlenecks	4
2. Using the Azure Portal to identify performance bottlenecks	9
3. Using TSQL and the Query Store to identify the root cause of the CPU strain	11

Migration architecture and Azure components



SQLHACK-SHARED-VNET		
Single Virtual Network containing all workshop resources		
TeamJumpServers Subnet Each team is assigned a Win10 VM that mimics their company desktop	Management Subnet Several machines and services are already deployed within a dedicated subnet within the Virtual Network	ManagedInstance Subnet The Azure SQL Managed Instance has been deployed into a dedicated Subnet

Lab Overview & Background

In this exercise you will explore how you can monitor and diagnose performance issues with Azure SQL Database Managed Instance using SQL Server Management Studio, DMVs (Dynamic Management Views) and Azure Portal tools.

Background

We have seen that the legacy application is a multi-tenant system with a collection of 3 databases supporting the transactional workload for each customer – these are the 3 databases we have modernised by migrating to the shared Azure SQL Managed Instance.

Apart from the tenant centric transactional databases there are 2 shared databases:

- **2008DW:** A centralised Data Warehouse database that combines data from the various tenant transactional databases and is used for aggregated reporting and analytics.
- **TenantCRM:** A centralised database that is used to manage all customer relationships and order processing.

Recently you have heard complaints from users that the TenantCRM system is running slowly. Recently changes were made to some of the stored procedures in the TenantCRM database and the App dev team believe these are the source of the problem. They've asked for your help to track down the performance issues and its root cause.

Let's follow the steps below to begin to troubleshoot this issue.

The exercises in this lab are conducted against the shared [TenantCRM] database

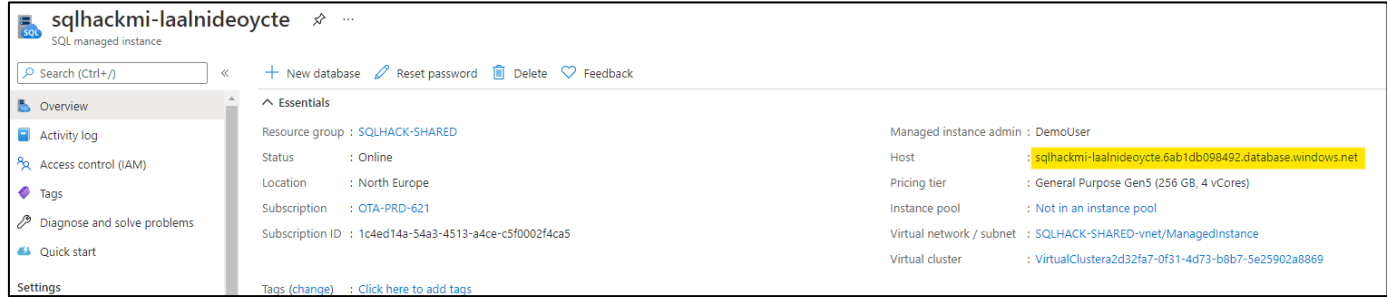
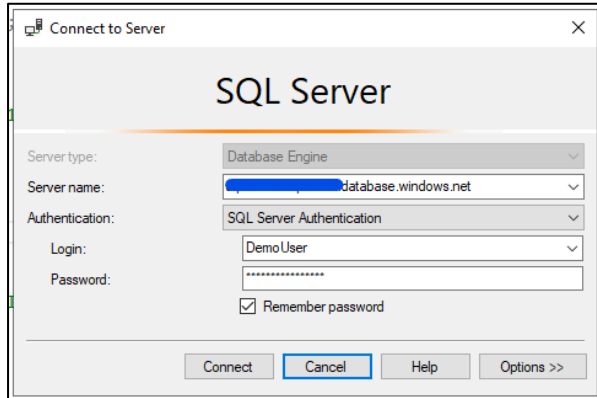
LAB 1: Identifying performance issues, their root causes and fixing the problem

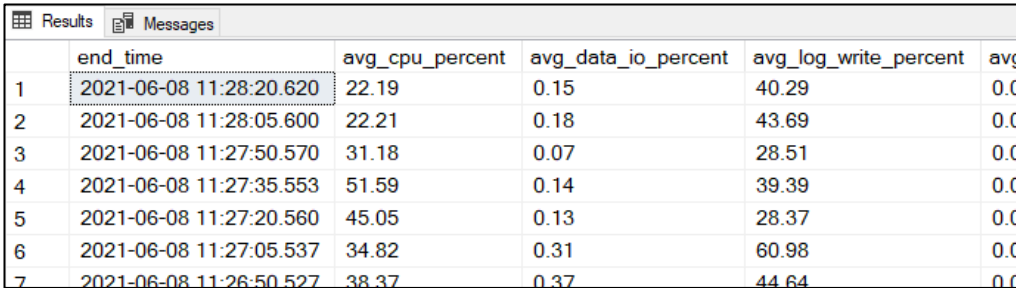
1. Using DMVs and the Query Store to identify performance bottlenecks.

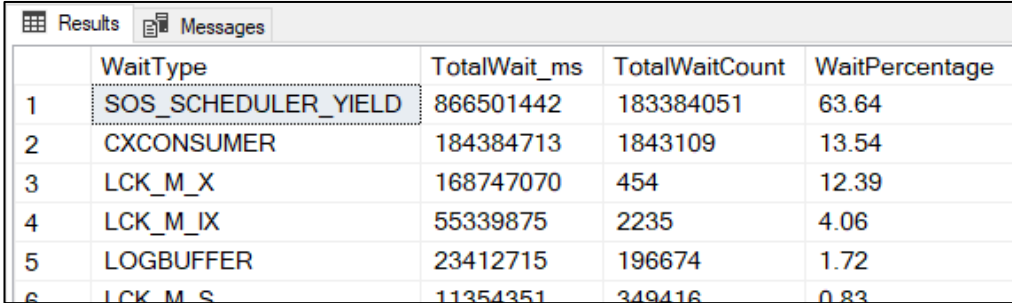
Performance issues in SQL Server can be grouped into 1 of 2 high-level categories: *Running* or *Waiting*.

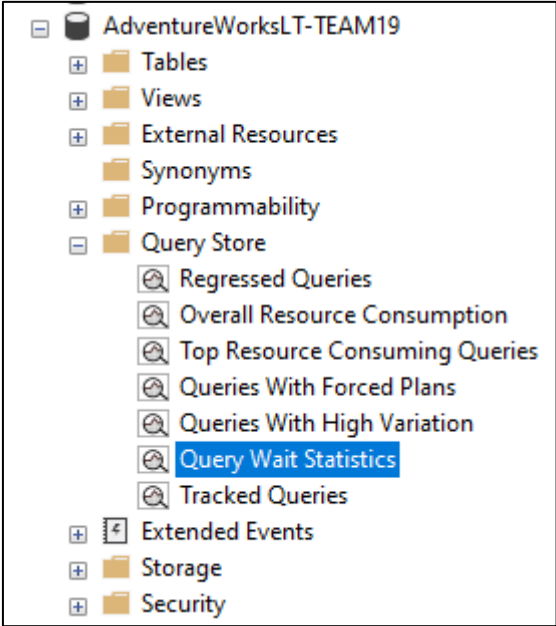
Running means a query has been allocated CPU resource and is actively processing but is taking a long time to complete. Waiting means a query is not actively being processed by the CPU but is waiting on another resource (such as memory, disk, network, etc) to pass data in so processing can continue.

Finding the bottleneck will help us determine what category issue we have so we can progress to identifying the root cause. Let's see how to find the bottleneck using both SSMS and the Azure portal.

Instructions/Narrative	Screenshot	Notes
<p>On your Win10 Team VM open the Azure Portal and navigate to the shared SQL Managed Instance.</p> <p>Scroll down the Overview screen copy the host name of the Managed Instance and save it somewhere handy.</p>		
<p>On your Team Win10 VM open SQL Server Management Studio and connect to the AdventureWorksLT using these details:</p> <p>Server name: <just copied> Authentication: SQL Server Login: DemoUser Pword: Demo@pass1234567</p>		

Narrative/Instructions	Screenshot	Notes																																																
<p>In SSMS open SQL script: C:_SQLHACK_\LABS\02-Administering_Monitoring\Part_02_Monitoring_Lab_1.sql connected to the shared [TenantCRM] database.</p> <p>Run the PART 1 query.</p> <p>Look at the results and note the [avg_cpu_percent] value is about 70% indicating that this database is consuming a lot of the hosts CPU resources</p> <p>(PTO)</p>	 <table><thead><tr><th></th><th>end_time</th><th>avg_cpu_percent</th><th>avg_data_io_percent</th><th>avg_log_write_percent</th><th>avg_memory_usage_percent</th></tr></thead><tbody><tr><td>1</td><td>2021-06-08 11:28:20.620</td><td>22.19</td><td>0.15</td><td>40.29</td><td>0.0</td></tr><tr><td>2</td><td>2021-06-08 11:28:05.600</td><td>22.21</td><td>0.18</td><td>43.69</td><td>0.0</td></tr><tr><td>3</td><td>2021-06-08 11:27:50.570</td><td>31.18</td><td>0.07</td><td>28.51</td><td>0.0</td></tr><tr><td>4</td><td>2021-06-08 11:27:35.553</td><td>51.59</td><td>0.14</td><td>39.39</td><td>0.0</td></tr><tr><td>5</td><td>2021-06-08 11:27:20.560</td><td>45.05</td><td>0.13</td><td>28.37</td><td>0.0</td></tr><tr><td>6</td><td>2021-06-08 11:27:05.537</td><td>34.82</td><td>0.31</td><td>60.98</td><td>0.0</td></tr><tr><td>7</td><td>2021-06-08 11:26:50.527</td><td>38.37</td><td>0.37</td><td>44.64</td><td>0.0</td></tr></tbody></table>		end_time	avg_cpu_percent	avg_data_io_percent	avg_log_write_percent	avg_memory_usage_percent	1	2021-06-08 11:28:20.620	22.19	0.15	40.29	0.0	2	2021-06-08 11:28:05.600	22.21	0.18	43.69	0.0	3	2021-06-08 11:27:50.570	31.18	0.07	28.51	0.0	4	2021-06-08 11:27:35.553	51.59	0.14	39.39	0.0	5	2021-06-08 11:27:20.560	45.05	0.13	28.37	0.0	6	2021-06-08 11:27:05.537	34.82	0.31	60.98	0.0	7	2021-06-08 11:26:50.527	38.37	0.37	44.64	0.0	<p>[sys].[dm_db_resource_stats] exposes a number of key database performance metrics including:</p> <p>avg_cpu_percent Average compute utilization in percentage of the limit of the service tier.</p> <p>avg_data_io_percent Average data I/O utilization as a percentage of the limit of the service tier</p> <p>avg_log_write_percent Average transaction log writes (in MBps) as percentage of the service tier limit.</p> <p>avg_memory_usage_percent Average memory utilization in percentage of the limit of the service tier.</p>
	end_time	avg_cpu_percent	avg_data_io_percent	avg_log_write_percent	avg_memory_usage_percent																																													
1	2021-06-08 11:28:20.620	22.19	0.15	40.29	0.0																																													
2	2021-06-08 11:28:05.600	22.21	0.18	43.69	0.0																																													
3	2021-06-08 11:27:50.570	31.18	0.07	28.51	0.0																																													
4	2021-06-08 11:27:35.553	51.59	0.14	39.39	0.0																																													
5	2021-06-08 11:27:20.560	45.05	0.13	28.37	0.0																																													
6	2021-06-08 11:27:05.537	34.82	0.31	60.98	0.0																																													
7	2021-06-08 11:26:50.527	38.37	0.37	44.64	0.0																																													

Narrative/Instructions	Screenshot	Notes																																			
<p>Let’s analyse the wait stats as you would in an on-premises environment.</p> <p>In SSMS run PART 2 of Part_04_Monitoring_Lab_1.sql which uses [sys].[dm_db_wait_stats] to report on top wait reasons since the instance was last restarted.</p> <p>View the results. You should see that SOS_SCHEDULER_YIELD is one of the top waits. This wait is reported when there are not enough CPU worker threads to fulfil the query demands.</p> <p>This confirms the results of [sys].[dm_db_resource_stats] DMV in the first query – that this database is putting the CPU under pressure.</p> <p>Now let’s see if the Query Store also points to CPU being the issue.</p>	 <table><thead><tr><th></th><th>WaitType</th><th>TotalWait_ms</th><th>TotalWaitCount</th><th>WaitPercentage</th></tr></thead><tbody><tr><td>1</td><td>SOS_SCHEDULER_YIELD</td><td>866501442</td><td>183384051</td><td>63.64</td></tr><tr><td>2</td><td>CXCONSUMER</td><td>184384713</td><td>1843109</td><td>13.54</td></tr><tr><td>3</td><td>LCK_M_X</td><td>168747070</td><td>454</td><td>12.39</td></tr><tr><td>4</td><td>LCK_M_IX</td><td>55339875</td><td>2235</td><td>4.06</td></tr><tr><td>5</td><td>LOGBUFFER</td><td>23412715</td><td>196674</td><td>1.72</td></tr><tr><td>6</td><td>LCK_M_S</td><td>11354351</td><td>349416</td><td>0.83</td></tr></tbody></table>		WaitType	TotalWait_ms	TotalWaitCount	WaitPercentage	1	SOS_SCHEDULER_YIELD	866501442	183384051	63.64	2	CXCONSUMER	184384713	1843109	13.54	3	LCK_M_X	168747070	454	12.39	4	LCK_M_IX	55339875	2235	4.06	5	LOGBUFFER	23412715	196674	1.72	6	LCK_M_S	11354351	349416	0.83	<p>[sys].[dm_db_resource_stats] exposes telemetry captured in the Query Store which you may have used in your on-premise environments. In Azure it is enabled for all Azure SQL Managed Instance and Azure SQL Database databases.</p> <p>The Query Store collects telemetry data every 15 seconds and persists it for about 1hr.</p> <p><i>The query used is taken from Glenn Berry's excellent SQL Server performance focused website https://glennsqlperformance.com/</i></p> <p><i>Glenn's website has a host of SQL version specific queries and is a great resource if you want to really stretch your DMV usage.</i></p>
	WaitType	TotalWait_ms	TotalWaitCount	WaitPercentage																																	
1	SOS_SCHEDULER_YIELD	866501442	183384051	63.64																																	
2	CXCONSUMER	184384713	1843109	13.54																																	
3	LCK_M_X	168747070	454	12.39																																	
4	LCK_M_IX	55339875	2235	4.06																																	
5	LOGBUFFER	23412715	196674	1.72																																	
6	LCK_M_S	11354351	349416	0.83																																	

Instructions/Narrative	Screenshot	Notes
<p>In SSMS expand the [TenantCRM] database then expand Query Store.</p> <p>Remember that the Query Store collects telemetry data every 15 seconds and persists it for about 1hr.</p> <p>Double click on Query Wait Statistics report to open it.</p> <p>(PTO)</p>	 <p>The screenshot shows the SQL Server Enterprise Manager interface. The 'AdventureWorksLT-TEAM19' database is selected and expanded. Under the 'Query Store' folder, the 'Query Wait Statistics' report is highlighted with a blue selection box. Other items visible in the tree include Tables, Views, External Resources, Synonyms, Programmability, Regressed Queries, Overall Resource Consumption, Top Resource Consuming Queries, Queries With Forced Plans, Queries With High Variation, Tracked Queries, Extended Events, Storage, and Security.</p>	

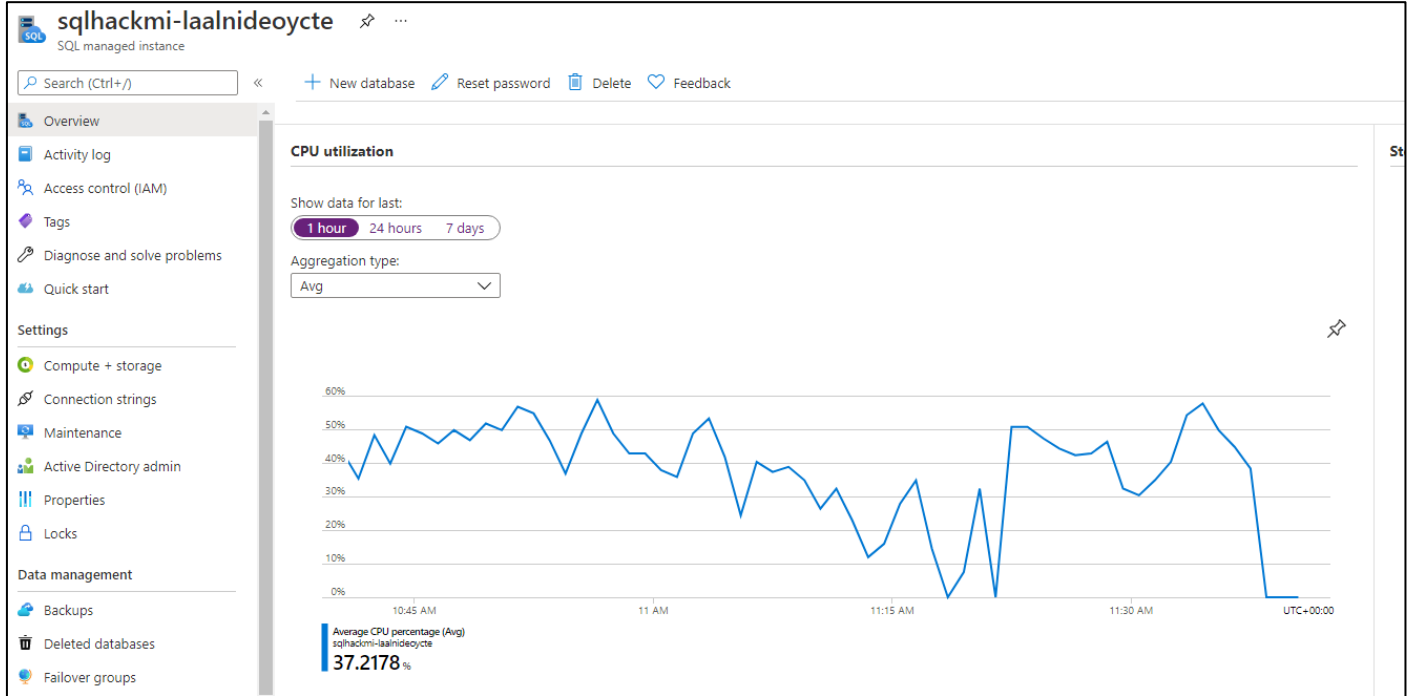
Narrative/Instructions	Screenshot	Notes																						
<p>You can see that CPU is also reflected here as one of the top waits on the database.</p> <p>Given that SOS_SCHEDULER_YIELD waits are CPU waits the Query Store report confirms CPU is [TenantCRM] database is dominating the CPU resources.</p>	<p>Query wait statistics for database AdventureWorksLT-TEAM19. Time period: Last hour ending at 4/3/2021 3:52 PM</p> <p>Based on: total wait time</p> <table><thead><tr><th>wait category</th><th>total wait time</th></tr></thead><tbody><tr><td>CPU</td><td>6800000</td></tr><tr><td>Latch</td><td>2200000</td></tr><tr><td>Memory</td><td>500000</td></tr><tr><td>Unknown</td><td>400000</td></tr><tr><td>Parallelism</td><td>300000</td></tr><tr><td>Log Rate Governor</td><td>100000</td></tr><tr><td>Lock</td><td>50000</td></tr><tr><td>Buffer Latch</td><td>20000</td></tr><tr><td>Tran Log IO</td><td>10000</td></tr><tr><td>Buffer IO</td><td>5000</td></tr></tbody></table>	wait category	total wait time	CPU	6800000	Latch	2200000	Memory	500000	Unknown	400000	Parallelism	300000	Log Rate Governor	100000	Lock	50000	Buffer Latch	20000	Tran Log IO	10000	Buffer IO	5000	
wait category	total wait time																							
CPU	6800000																							
Latch	2200000																							
Memory	500000																							
Unknown	400000																							
Parallelism	300000																							
Log Rate Governor	100000																							
Lock	50000																							
Buffer Latch	20000																							
Tran Log IO	10000																							
Buffer IO	5000																							

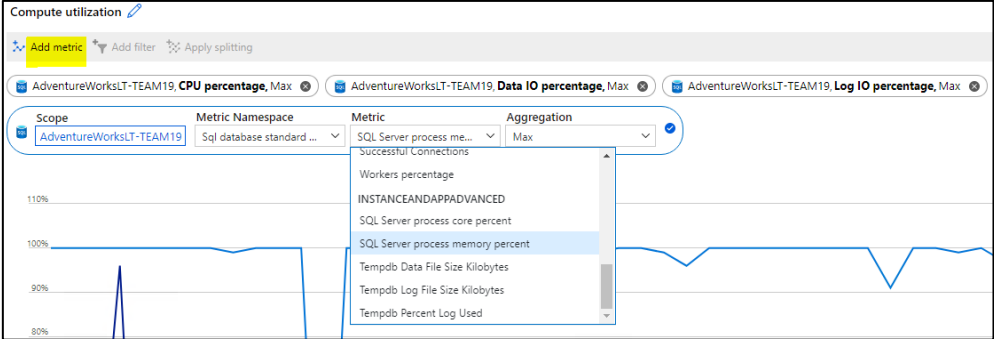
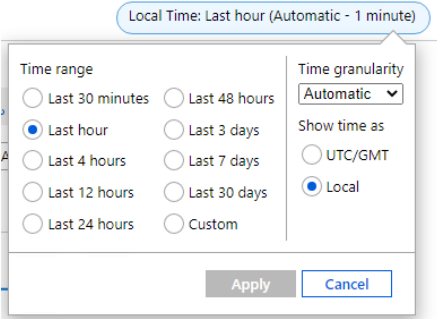
(PTO)

SQL Modernisation Open Hack

2. Using the Azure Portal to identify performance bottlenecks

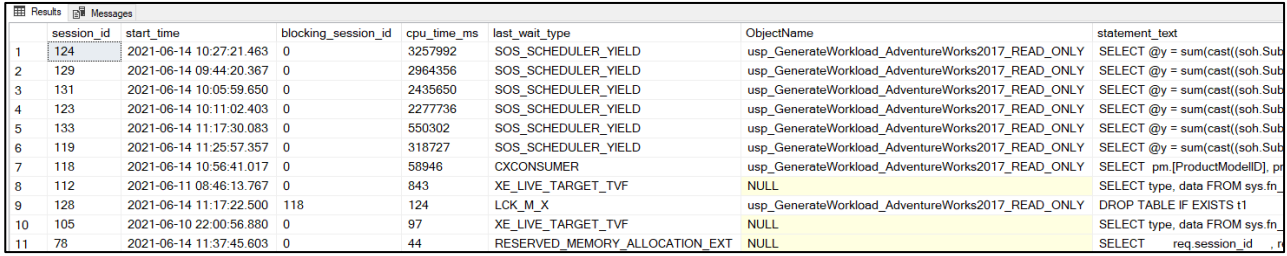
As we are using Azure SQL Managed Instance we can also use the richness of the Azure Portal to help us to troubleshoot performance issues. Let us see how this is done.

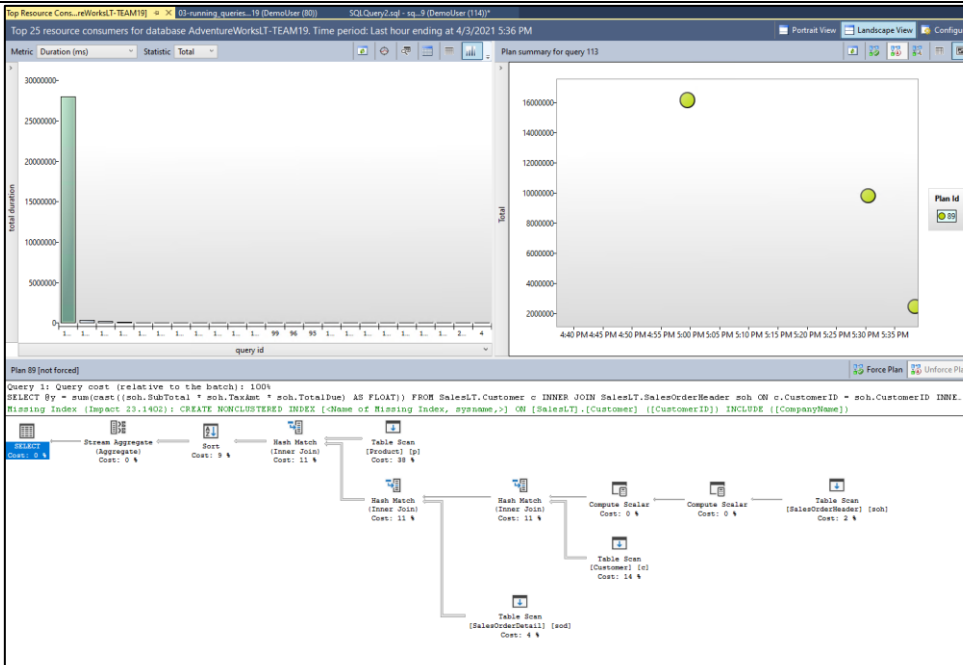
Instructions/Narrative	Screenshot	Notes
<p>Open the Azure Portal (portal.azure.com) and open the shared SQL Managed Instance.</p> <p>Scroll down the Overview screen until you see the performance graph showing CPU utilization.</p> <p>Toggle the switches next to view CPU consumption over 1 hour, 24 hours and 7 days.</p> <p>Has the CPU usage pattern changed over time?</p>	 <p>The screenshot shows the Azure Portal interface for a SQL Managed Instance named 'sqlhackmi-laalnideoycte'. The left sidebar contains navigation options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Quick start, Settings, Compute + storage, Connection strings, Maintenance, Active Directory admin, Properties, Locks, Data management, Backups, Deleted databases, and Failover groups. The main area displays the 'CPU utilization' graph. The graph shows a line representing CPU usage over time, with a peak around 55% and a current average of 37.2178%. The x-axis shows time from 10:45 AM to 11:30 AM, and the y-axis shows CPU percentage from 0% to 60%.</p>	

Instructions/Narrative	Screenshot	Notes
<p>Click into the graph to go into Metrics chart screen. Here you have more flexibility over the data being reported and the format of the chart. Let's add IO Request Count the chart.</p> <p>Click Add Metric</p> <p>In the Metric drop down select IO requests count.</p> <p>Note that because the Average CPU Percentage metric was already on the chart the vertical axis uses percentage which doesn't make much sense when we're looking at the count of IO requests. We will fix this in a later lab.</p>		<p>Notice that from the chart screen you can alter the type of chart (Line, Area, Bar, Scatter to data Grid), set up alerts, and pin the chart to your Azure Portal dashboard.</p>
<p>Click Local Time on the top right of the chart to view the data for various time intervals.</p>		

3. Using TSQL and the Query Store to identify the root cause of the CPU strain

Now that we have discovered that the amount of CPU consumed by the [TenantCRM] database is a concern, we need to drill-down and find out what batches are causing the issue. Again we can do this through both SSMS and the Azure portal. Let us begin with Management Studio.

Instructions/Narrative	Screenshot	Notes
<p>In SSMS return to the SQL script: C:_SQLHACK_\LABS\02-Administering_Monitoring\Part_04_Monitoring_Lab_1.sql which should still be connected to your connected to the [TenantCRM] database.</p> <p>Run the PART 3 query which uses the sys.dm_exec_requests and sys.dm_exec_sql_text DMVs.</p> <p>View the results. Note columns last_wait_type and statement_text as statement_text reveals the poorly running query.</p> <p>We have already identified that SOS_SCHEDULER_YIELD is the wait type associated with CPU stress and we can now see the query which is causing this bottleneck to occur.</p> <p>Run this query a few times to see if the suspect query and wait type is consistent over multiple runs.</p>		
	<p>CROSS APPLY and Table-valued Functions</p> <p>The PART 3 query joins the [sys].[dm_exec_requests] and [sys].[dm_exec_sql_text] DMVs to obtain long running batches and critically their offending SQL.</p> <p>Note that [dm_exec_sql_text] is actually a table-valued function (TVF) hence the use of the CROSS APPLY as TVFs can't be used in a normal join operation. The CROSS APPLY therefore produces the required inner-join between [dm_exec_requests] and [dm_exec_sql_text].</p> <p>See APPLY documentation here: Using APPLY Microsoft Docs and this excellent explanation on MSSQLTips: SQL Server CROSS APPLY and OUTER APPLY (mssqltips.com)</p>	

Instructions/Narrative	Screenshot	Notes
<p>In SSMS expand your database and open the Query Store.</p> <p>Double click Top Resource Consuming Queries report to open it.</p> <p>Click on the tallest bar which represents the query with the longest duration.</p> <p>The query text and execution plan will be shown below the charts allowing us to identify the offending query text. Note that it is the same as you saw in the previous DMV output.</p>	 <p>The screenshot displays the SQL Server Query Store interface. At the top, the 'Top Resource Consuming Queries' report is open, showing a bar chart of total duration for various queries. The tallest bar is selected, and the query text and execution plan are displayed below. The query is a SELECT statement with a subquery and a JOIN. The execution plan shows a complex join structure with multiple table scans and joins.</p> <p>Query 1: Query cost (relative to the batch): 100%</p> <p>SELECT By = sum(cost*((sub.SubTotal * sub.Taxamt * sub.TotalDue) AS FLOAT)) FROM SalesLT.Customer c INNER JOIN SalesLT.SalesOrderHeader soh ON c.CustomerID = soh.CustomerID INNER JOIN SalesLT.Product p ON soh.ProductID = p.ProductID WHERE (c.CompanyName = 'AdventureWorks') AND (p.ProductID = 1)</p> <p>Missing Index (Impact: 22.140%): CREATE NONCLUSTERED INDEX (Name of Missing Index, sysname...) ON (SalesLT].[Customer]) ((CustomerID)) INCLUDE ((CompanyName))</p> <p>Execution Plan: The plan shows a complex join structure. It starts with a Table Scan (SalesLT.Product) [p] (Cost: 28 %), followed by a Hash Match (Inner Join) (Cost: 11 %), then a Stream Aggregate (Aggregate) (Cost: 0 %), and finally a Table Scan (SalesLT.Customer) [c] (Cost: 14 %). The plan also includes a Table Scan (SalesOrderHeader) [soh] (Cost: 2 %).</p>	<p>The graph on the right shows all the different execution plans for the same query (there will probably only be 1 plan) and their execution duration over time.</p>

We've now seen how the various performance monitoring and diagnosis tools – the Azure Portal, DMVs and the Query Store reports all revealed that the CPU was under pressure. The DMVs and Query Store reports also allowed us to drill-down to root cause and identify the worst offending query. In the real world you would now progress to tune the offending query to reduce the overall load on the Managed Instances CPUs and thereby improve the databases and environment performance.